

Figure 7. Δ judge accuracy by question type $\times$ readiness dimension (best epoch per model: Qwen3-8B ep1, Gemma 4 E4B ep1, Llama 3.1 8B ep2; all splits pooled)

