

judge accuracy, mean of 3 tests

- Qwen3-8B
- Gemma 4 E4B
- Llama 3.1 8B
- bf16 base + LoRA (vLLM)
- NF4 base + LoRA (transformers)
- merged 4-bit (vLLM)

